

# Architecture Decision Record

*Git CMS*

---

Database-Free Content Management via Git Plumbing

**Author:** James Ross

**Version:** 1.0.0

**Date:** 2026-01-11

*Prepared for Engineering Review*

## Contents

<b>1</b>	<b>Introduction &amp; Goals</b>	<b>2</b>
1.1	Project Overview . . . . .	2
1.2	Fundamental Requirements . . . . .	2
1.2.1	FR-1: Zero-Database Architecture . . . . .	2
1.2.2	FR-2: Cryptographic Verifiability . . . . .	2
1.2.3	FR-3: Fast-Forward Only Publishing . . . . .	2
1.2.4	FR-4: Client-Side Encryption . . . . .	2
1.2.5	FR-5: Infinite Point-in-Time Recovery . . . . .	2
1.3	Quality Goals . . . . .	3
1.4	Non-Goals . . . . .	3
<b>2</b>	<b>Constraints</b>	<b>3</b>
2.1	Technical Constraints . . . . .	3
2.2	Regulatory Constraints . . . . .	4
2.3	Operational Constraints . . . . .	4
<b>3</b>	<b>Context &amp; Scope</b>	<b>4</b>
3.1	System Context Diagram . . . . .	4
3.2	External Interfaces . . . . .	5
3.2.1	Interface 1: CLI (Binary) . . . . .	5
3.2.2	Interface 2: HTTP API (REST) . . . . .	5
3.2.3	Interface 3: Git Plumbing (Shell) . . . . .	5
3.2.4	Interface 4: OS Keychain (Secrets) . . . . .	5
3.3	Scope Boundaries . . . . .	5
3.3.1	In Scope . . . . .	5
3.3.2	Out of Scope . . . . .	5
<b>4</b>	<b>Solution Strategy</b>	<b>6</b>
4.1	Core Architectural Principles . . . . .	6
4.2	Solution Approach: The Empty TreeStunt . . . . .	6
4.3	Key Design Decisions . . . . .	7
<b>5</b>	<b>Building Block View</b>	<b>8</b>
5.1	Level 1: System Decomposition . . . . .	8
5.2	Level 2: Lego Block Responsibilities . . . . .	9
<b>6</b>	<b>Runtime View</b>	<b>10</b>
6.1	Scenario 1: Create Draft Article . . . . .	10
6.2	Scenario 2: Publish Article . . . . .	10
6.3	Scenario 3: Upload Encrypted Asset . . . . .	10
6.4	Scenario 4: List All Published Articles . . . . .	10
<b>7</b>	<b>Deployment View</b>	<b>11</b>
7.1	Topology 1: Single-Author Local Blog . . . . .	11
7.2	Topology 2: Team Blog with Stargate Gateway . . . . .	11
7.3	Topology 3: Dockerized Development . . . . .	12
<b>8</b>	<b>Crosscutting Concepts</b>	<b>13</b>
8.1	Concept 1: Merkle DAG as Event Log . . . . .	13
8.2	Concept 2: Compare-and-Swap (CAS) . . . . .	13

8.3 Concept 3: Client-Side Encryption . . . . .	14
<b>9 Architectural Decisions</b>	<b>14</b>
<b>10 Quality Requirements</b>	<b>15</b>
10.1 Quality Tree . . . . .	15
10.2 Quality Scenarios . . . . .	15
10.2.1 QS-1: Tamper Detection . . . . .	15
10.2.2 QS-2: Confidentiality . . . . .	16
<b>11 Risks &amp; Technical Debt</b>	<b>16</b>
11.1 Risk 1: SHA-1 Collision . . . . .	16
11.2 Risk 2: Repository Growth . . . . .	16
11.3 Technical Debt Summary . . . . .	16
<b>12 Glossary</b>	<b>16</b>
<b>A Appendix A: Example Commands</b>	<b>16</b>
A.1 Draft an Article . . . . .	16
A.2 Publish an Article . . . . .	17
A.3 List All Drafts . . . . .	17
A.4 Upload Asset . . . . .	17
<b>B Appendix B: Directory Structure</b>	<b>17</b>
<b>C Appendix C: Related Projects</b>	<b>17</b>
<b>D Appendix D: References</b>	<b>17</b>
<b>E Appendix E: References</b>	<b>17</b>

# 1 Introduction & Goals

## 1.1 Project Overview

**git-cms** is a serverless, database-free Content Management System that treats Git's object store as a distributed, cryptographically verifiable document database. Instead of storing content in traditional databases (SQL or NoSQL), it leverages Git's Merkle DAG to create an append-only ledger for articles, metadata, and encrypted assets.

The fundamental innovation: `git push` becomes the API endpoint.

## 1.2 Fundamental Requirements

### 1.2.1 FR-1: Zero-Database Architecture

The system MUST NOT depend on external database systems (SQL, NoSQL, or key-value stores). All persistent state resides within Git's native object store (`.git/objects`).

**Rationale:** Eliminates operational complexity, deployment dependencies, and schema migration challenges inherent to traditional database-backed CMSs.

### 1.2.2 FR-2: Cryptographic Verifiability

Every content mutation MUST be recorded as a Git commit with cryptographic integrity guarantees via SHA-1 hashing (with optional GPG signing for non-repudiation).

**Rationale:** Provides immutable audit trails and tamper detection without additional infrastructure.

### 1.2.3 FR-3: Fast-Forward Only Publishing

The publish operation MUST enforce strict linear history (fast-forward only) to prevent rewriting published content.

**Rationale:** Guarantees provenance and prevents content manipulation after publication.

### 1.2.4 FR-4: Client-Side Encryption

All uploaded assets MUST be encrypted client-side (AES-256-GCM) before touching the repository.

**Rationale:** Achieves row-level security without database-level access controls. The Git gateway receives only opaque encrypted blobs.

### 1.2.5 FR-5: Infinite Point-in-Time Recovery

Users MUST be able to access any historical version of any article without data loss.

**Rationale:** Git's DAG structure provides this naturally; the CMS simply exposes it as a first-class feature.

### 1.3 Quality Goals

Prio	Attribute	Description	Measurement
1	Security	Cryptographic integrity, signed commits	PGP verification, AES-256 strength
2	Simplicity	Minimal dependencies, composable architecture	Lines of code, dependency count
3	Auditability	Complete provenance of all content changes	Git log completeness
4	Performance	Sub-second reads for blog workloads	Response time for <code>readArticle()</code>
5	Portability	Multi-runtime support (Node, Bun, Deno)	Test suite pass rate

Table 1: Quality goals and their measurements.

### 1.4 Non-Goals

This system is **intentionally NOT designed for:**

- **High-velocity writes:** Content publishing happens in minutes/hours, not milliseconds.
- **Complex queries:** No SQL-like JOINs or aggregations. Queries are limited to ref enumeration and commit message parsing.
- **Large-scale collaboration:** Designed for single-author or small-team blogs.
- **Real-time updates:** Publishing is atomic but not instantaneous.

## 2 Constraints

### 2.1 Technical Constraints

#### TC-1: Git's Content Addressability Model

Git uses SHA-1 hashing for object addressing. While SHA-1 has known collision vulnerabilities, Git is transitioning to SHA-256. The system assumes SHA-1 is “good enough” for content addressing (not for security-critical signing).

**Mitigation:** Use GPG signing (`CMS_SIGN=1`) for cryptographic non-repudiation.

#### TC-2: Filesystem I/O Performance

All Git operations are ultimately filesystem operations. Performance is bounded by disk I/O, especially for large repositories.

**Mitigation:** Content is stored as commit messages (small), not files (large). Asset chunking (256KB) reduces blob size.

#### TC-3: POSIX Shell Dependency

The `@git-stunts/plumbing` module executes Git via shell commands (`child_process.spawn`). This requires a POSIX-compliant shell and Git CLI.

**Mitigation:** All tests run in Docker (Alpine Linux) to ensure consistent environments.

#### TC-4: No Database Indexes

Traditional databases provide B-tree indexes for fast lookups. Git’s ref enumeration is linear ( $O(n)$  for listing all refs in a namespace).

**Mitigation:** Use ref namespaces strategically (e.g., `refs/_blog/articles/<slug>`) to avoid

polluting the global ref space.

## 2.2 Regulatory Constraints

### RC-1: GDPR Right to Erasure

Git's immutability conflicts with GDPR's "right to be forgotten." Deleting a commit requires rewriting history, which breaks cryptographic integrity.

**Mitigation:** Use encrypted assets with key rotation. Deleting the encryption key renders historical content unreadable without altering Git history.

### RC-2: Cryptographic Export Restrictions

AES-256-GCM encryption may face export restrictions in certain jurisdictions.

**Mitigation:** The `@git-stunts/vault` module uses Node's built-in `crypto` module, which is widely available.

## 2.3 Operational Constraints

### OC-1: Single-Writer Assumption

Git's ref updates are atomic *locally* but not across distributed clones. Concurrent writes to the same ref can cause conflicts.

**Mitigation:** Use `git-stargate` (a companion project) to enforce serialized writes via SSH.

### OC-2: Repository Growth

Every draft save creates a new commit. Repositories can grow unbounded over time.

**Mitigation:** Use `git gc` aggressively. Consider ref pruning for old drafts.

## 3 Context & Scope

### 3.1 System Context Diagram

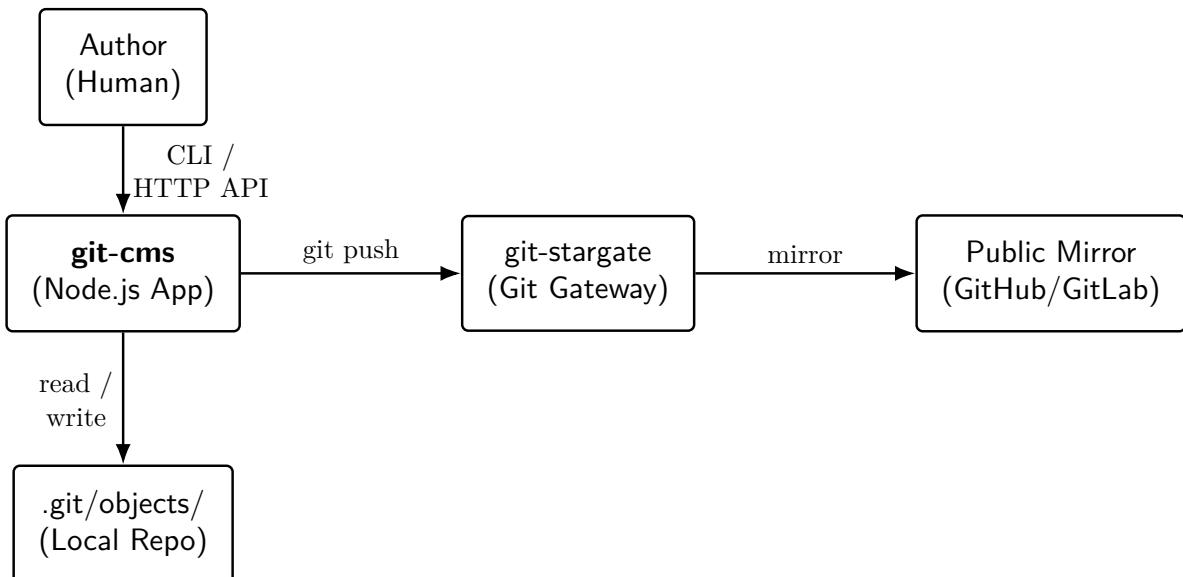


Figure 1: System context diagram showing the high-level relationship between the Author, Git CMS, and external components.

## 3.2 External Interfaces

### 3.2.1 Interface 1: CLI (Binary)

- **Entry Point:** bin/git-cms.js
- **Commands:** draft, publish, list, show, serve
- **Protocol:** POSIX command-line arguments
- **Example:**

```
1 echo "# Hello World" | git cms draft hello-world "My First Post"
```

### 3.2.2 Interface 2: HTTP API (REST)

- **Server:** src/server/index.js
- **Port:** 4638 (configurable via PORT env var)
- **Endpoints:**
  - POST /api/cms/snapshot – Save draft
  - POST /api/cms/publish – Publish article
  - GET /api/cms/list – List articles
  - GET /api/cms/show?slug=<slug> – Read article
- **Authentication:** None (assumes private network or SSH tunneling)

### 3.2.3 Interface 3: Git Plumbing (Shell)

- **Protocol:** Git CLI commands via child\_process.spawn
- **Critical Commands:**
  - git commit-tree – Create commits on empty trees
  - git update-ref – Atomic ref updates
  - git for-each-ref – List refs in namespace
  - git cat-file – Read commit messages

### 3.2.4 Interface 4: OS Keychain (Secrets)

- **Platforms:**
  - macOS: security tool
  - Linux: secret-tool (GNOME Keyring)
  - Windows: CredentialManager (PowerShell)
- **Purpose:** Store AES-256-GCM encryption keys for assets

## 3.3 Scope Boundaries

### 3.3.1 In Scope

- Article drafting, editing, and publishing
- Encrypted asset storage (images, PDFs)
- Full version history via Git log
- CLI and HTTP API access
- Multi-runtime support (Node, Bun, Deno)

### 3.3.2 Out of Scope

- **User Authentication:** Delegated to git-stargate or SSH
- **Search Indexing:** No full-text search (external indexer required)
- **Media Transcoding:** Assets stored as-is

- **Real-Time Collaboration:** No OT or CRDTs
- **Analytics:** No built-in tracking

## 4 Solution Strategy

### 4.1 Core Architectural Principles

**P-1: Composition over Inheritance** The system is built from **five independent Lego Blockmodules** (`@git-stunts/*`), each with a single responsibility. These modules are composed in `CmsService` to create higher-order functionality.

**Benefit:** Each module can be tested, versioned, and published independently.

**P-2: Hexagonal Architecture (Ports & Adapters)** The domain layer (`CmsService`) depends on abstractions (`GitPlumbing`, `TrailerCodec`), not implementations. This allows swapping out Git for other backends (e.g., a pure JavaScript implementation for testing).

**Benefit:** Decouples domain logic from infrastructure concerns.

**P-3: Content Addressability** Assets are stored by their SHA-1 hash, enabling automatic deduplication. If two articles reference the same image, it's stored once.

**Benefit:** Reduces repository bloat.

**P-4: Cryptographic Integrity** Every operation produces a cryptographically signed commit (when `CMS_SIGN=1`). The Merkle DAG ensures tamper detection.

**Benefit:** Audit trails are mathematically verifiable, not just trust-based.

### 4.2 Solution Approach: The Empty TreeStunt

**The Problem** Traditional CMSs store content in database rows. Git is designed to track *files*, not arbitrary data. Storing blog posts as files (e.g., `posts/hello-world.md`) clutters the working directory and causes merge conflicts.

**The Solution** Store content as **commit messages on empty trees**, not as files. Every article is a commit that points to the well-known empty tree (`4b825dc642cb6eb9a060e54bf8d69288fbee4904`).

#### How It Works:

1. Encode the article (title, body, metadata) into a Git commit message using RFC 822 trailers.
2. Create a commit that points to the empty tree (no files touched).
3. Update a ref (e.g., `refs/_blog/articles/hello-world`) to point to this commit.

**Result:** The repository's working directory remains clean. All content lives in `.git/objects/` and `.git/refs/`.

**Architectural Pattern: Event Sourcing** Each draft save creates a new commit. The current article is the ref's tip, but the full history is a linked list of commits.

**Benefit:** Point-in-time recovery is trivial (`git log refs/_blog/articles/<slug>`).

### 4.3 Key Design Decisions

**D-1: Why Commit Messages, Not Blobs?** **Alternative:** Store articles as Git blobs and reference them via trees.

**Decision:** Use commit messages.

**Rationale:** Commits have parent pointers (version history) and support GPG signing (non-repudiation). Blobs are opaque; messages are human-readable.

**D-2: Why Trailers, Not JSON?** **Alternative:** Store `{title: "Hello, ..."}` as the message.

**Decision:** Use RFC 822 trailers.

**Rationale:** Trailers are Git-native, human-readable, and diff-friendly. Backward parsing is efficient.

**D-3: Why Encrypt Assets, Not Repos?** **Alternative:** Use `git-crypt` for the whole repo.

**Decision:** Encrypt individual assets client-side.

**Rationale:** Granular control; the gateway never sees plaintext.

## 5 Building Block View

### 5.1 Level 1: System Decomposition

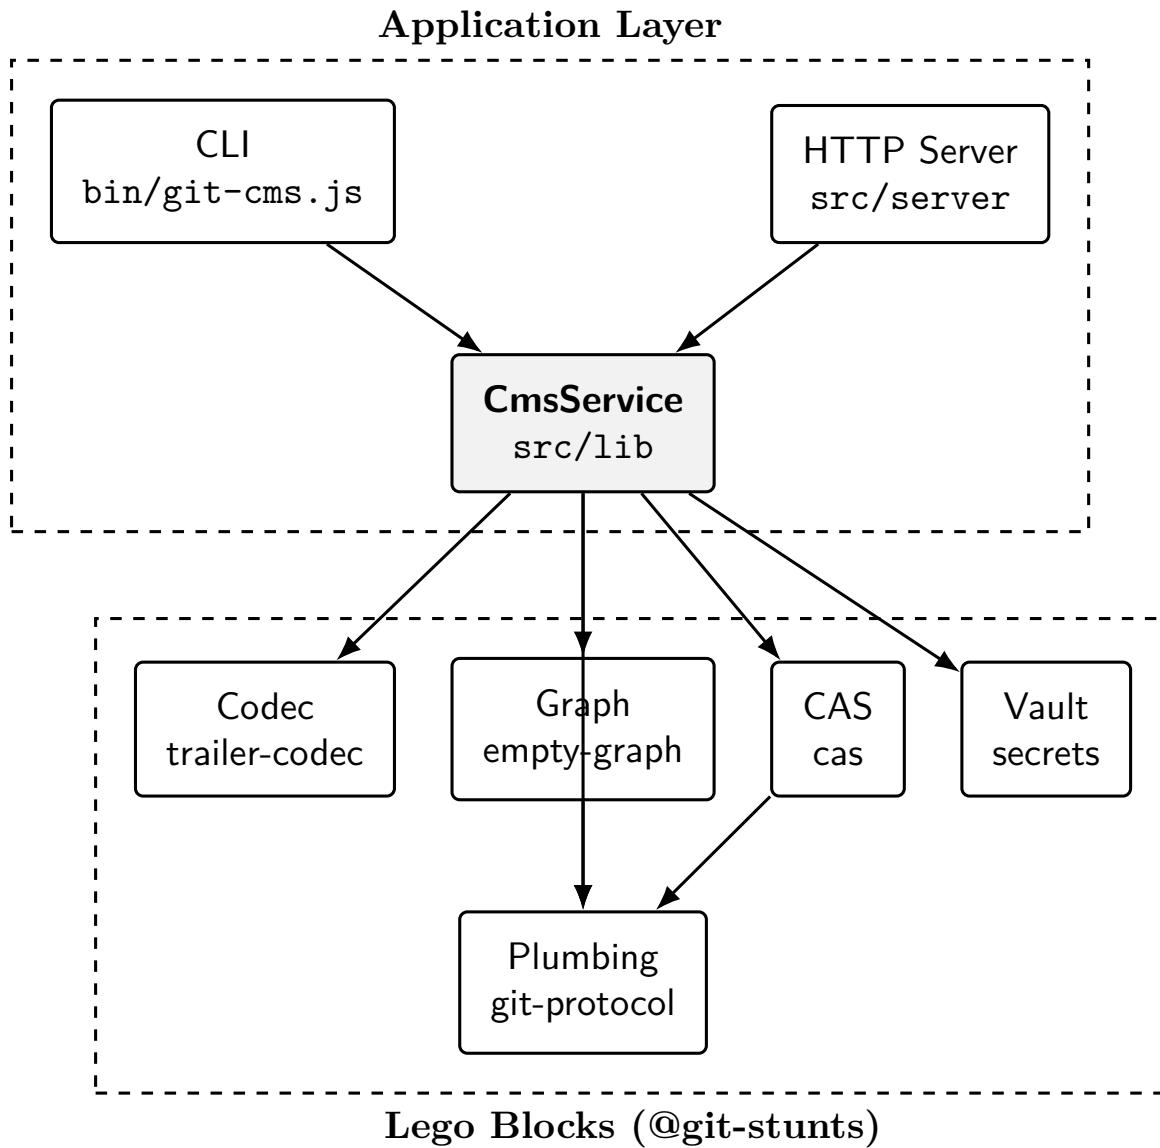


Figure 2: System decomposition showing the interaction between the application layer and the independent Lego block modules.

## 5.2 Level 2: Lego Block Responsibilities

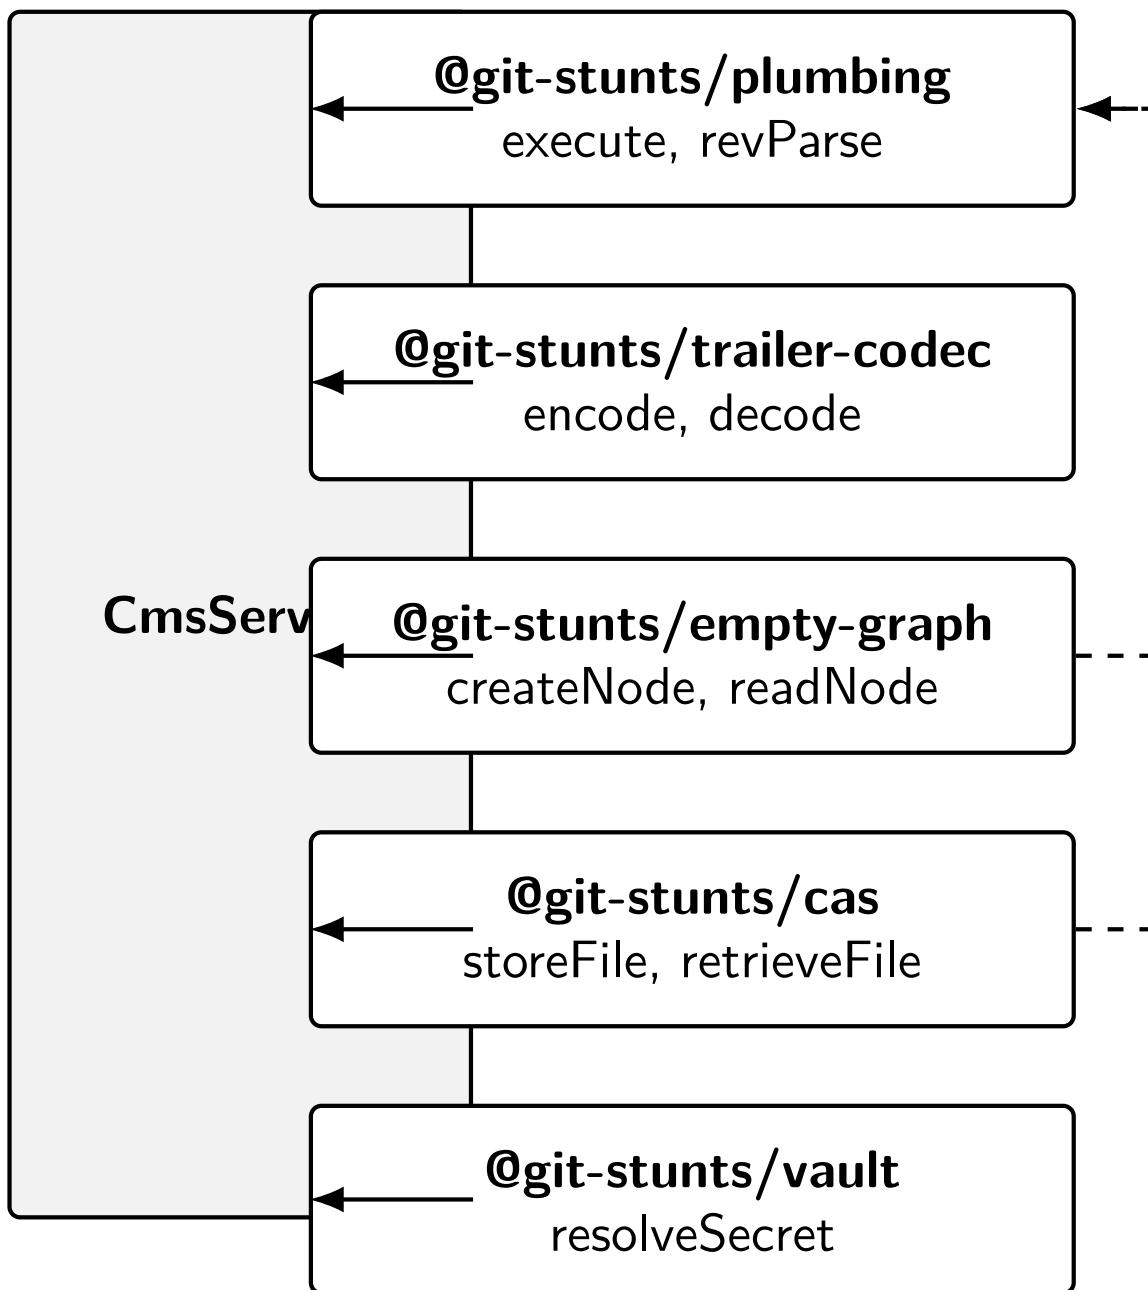


Figure 3: Detailed responsibilities and API surfaces of the Git CMS modules.

## 6 Runtime View

### 6.1 Scenario 1: Create Draft Article

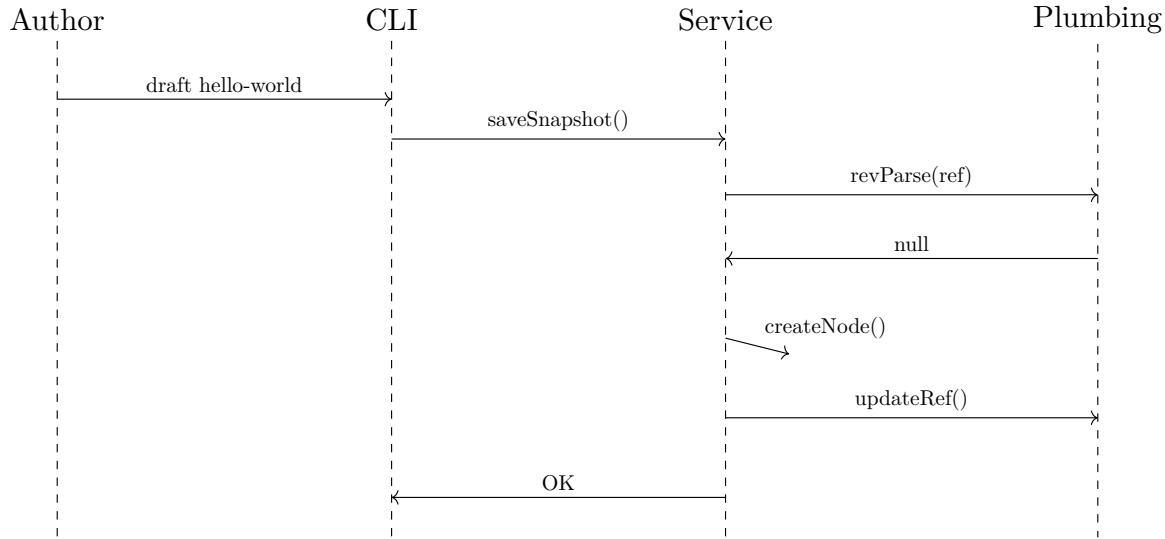


Figure 4: Sequence diagram for creating a new draft article.

### 6.2 Scenario 2: Publish Article

Publishing is **just a ref copy**. No new commits are created. This operation is idempotent and enforces fast-forward updates.

### 6.3 Scenario 3: Upload Encrypted Asset

The system splits files into 256KB chunks, encrypts them via AES-256-GCM, and stores them as Git blobs. The plaintext never touches the object store.

### 6.4 Scenario 4: List All Published Articles

Listing articles involves a linear scan of the ref namespace ( $O(n)$ ). For large workloads, an external index is recommended.

## 7 Deployment View

### 7.1 Topology 1: Single-Author Local Blog

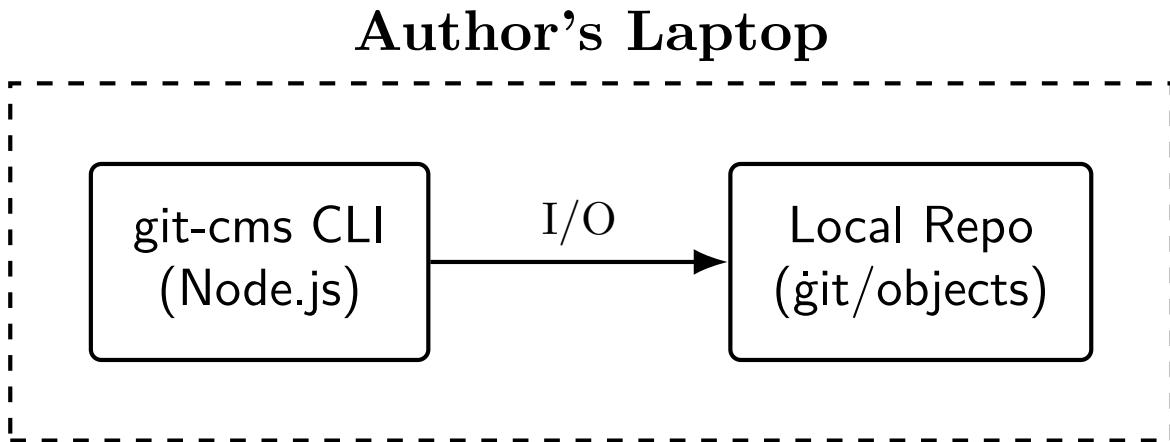


Figure 5: Local deployment topology.

### 7.2 Topology 2: Team Blog with Stargate Gateway

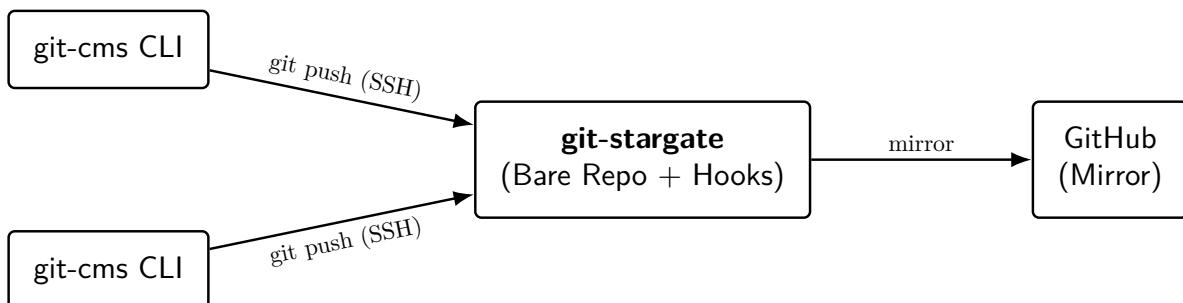


Figure 6: Collaborative deployment topology using a central gateway.

### 7.3 Topology 3: Dockerized Development

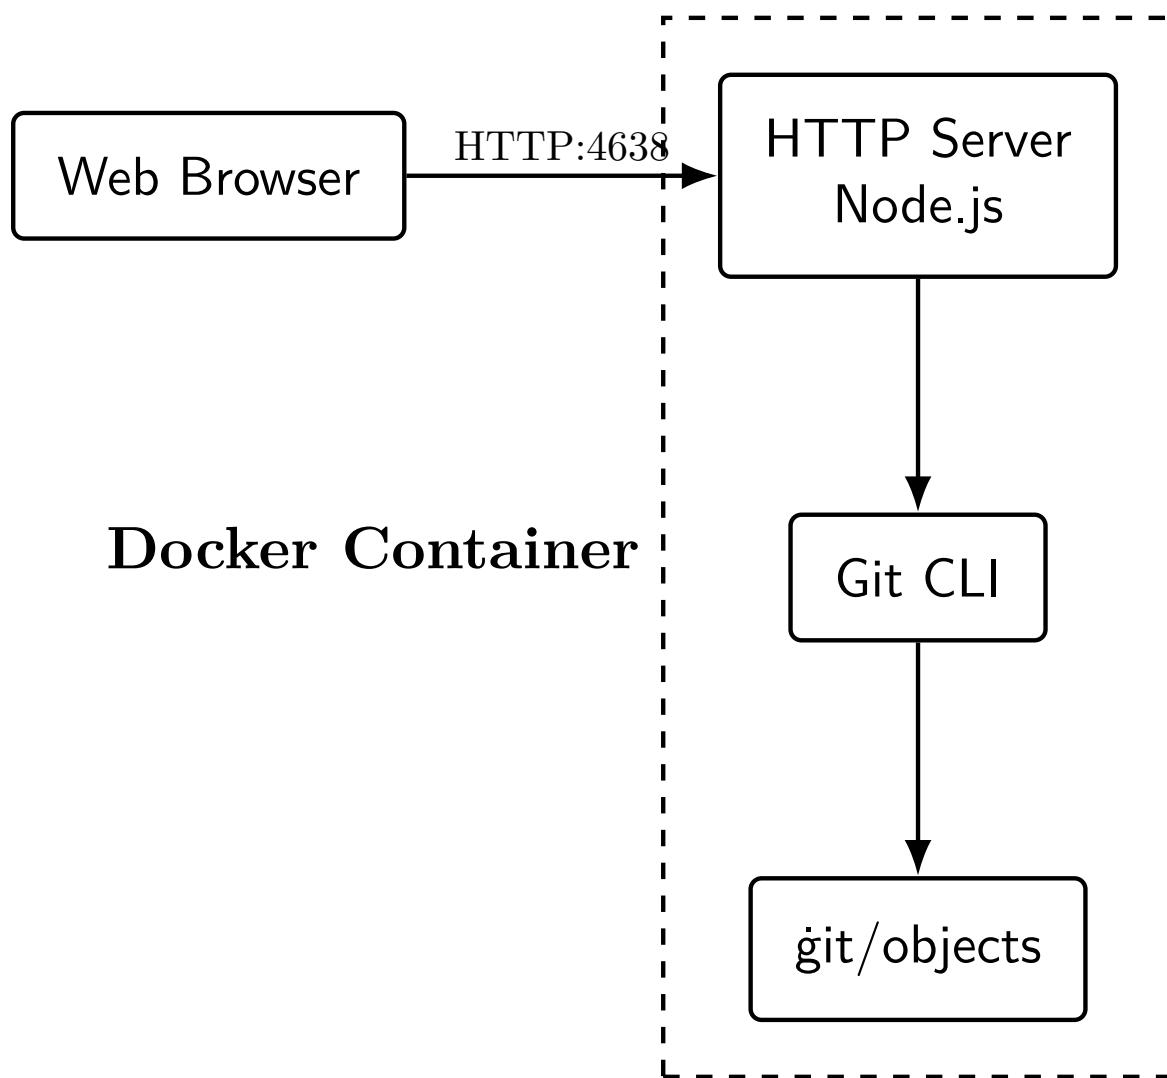


Figure 7: Dockerized development topology.

## 8 Crosscutting Concepts

### 8.1 Concept 1: Merkle DAG as Event Log

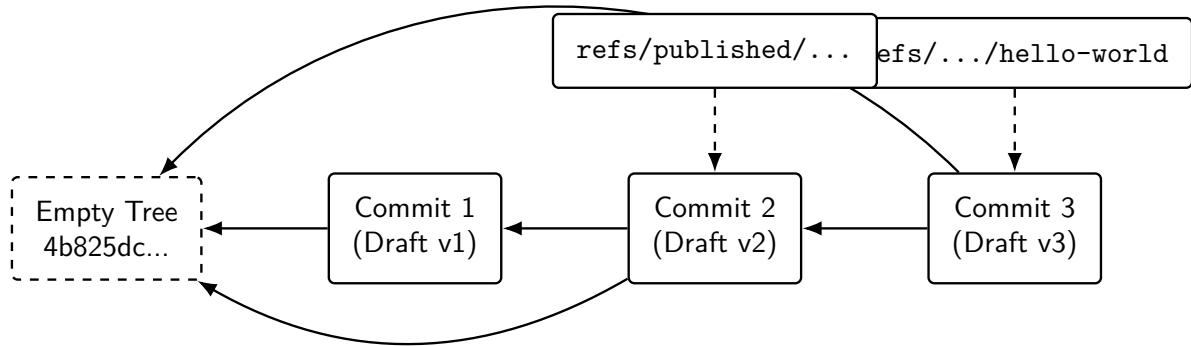


Figure 8: The Merkle DAG structure acting as an immutable event log.

### 8.2 Concept 2: Compare-and-Swap (CAS)

The system uses `git update-ref <ref> <newSHA> <oldSHA>` to ensure atomic updates and prevent race conditions. If the `oldSHA` has changed since it was last read, the update is rejected.

### 8.3 Concept 3: Client-Side Encryption

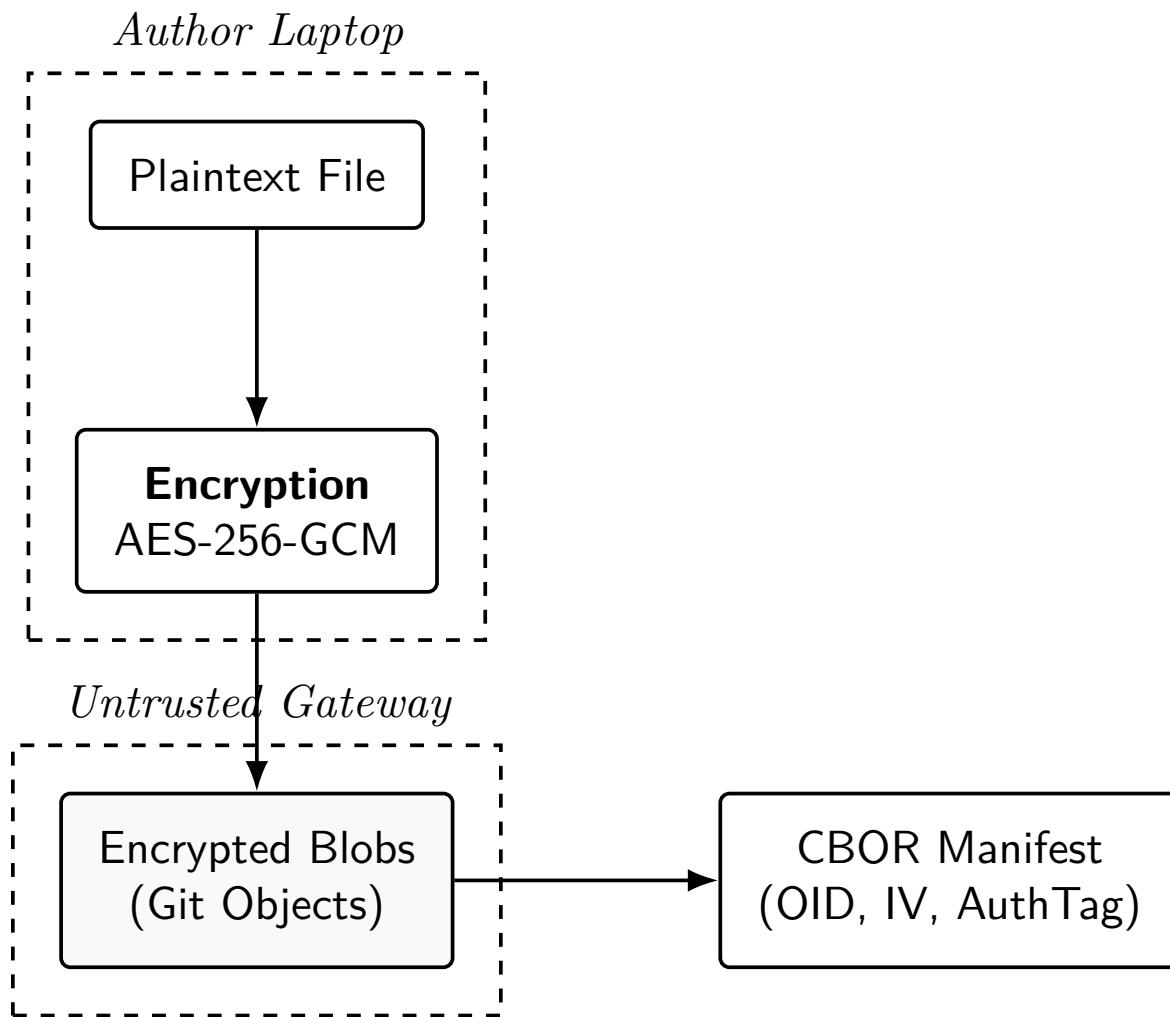


Figure 9: End-to-end encryption pipeline for assets.

## 9 Architectural Decisions

### ADR-001: Use Commit Messages, Not Files

**Context:** Need to store articles in Git without polluting the working directory or causing merge conflicts on files.

**Decision:** Store article content (title, body, trailers) as Git commit messages pointing to the canonical empty tree (4b825dc...).

#### Rationale:

- Commits have parent pointers, enabling native version history.
- Commits support GPG signing for non-repudiation.
- Keeps the working directory completely clean for application code.

**Status:** Accepted.

## ADR-002: Use RFC 822 Trailers, Not JSON

**Context:** Need structured metadata (Status, Author, etc.) inside commit messages.

**Decision:** Use RFC 822 trailers (key-value pairs at the end of the message).

**Rationale:**

- Git-native format (compatible with `git interpret-trailers`).
- Human-readable and extremely diff-friendly.
- Faster to parse from the end of the message.

**Status:** Accepted.

## ADR-003: Fast-Forward Only Publishing

**Context:** Prevent published content from being altered or rewritten after release.

**Decision:** The publishing operation must be a strict fast-forward from the draft ref to the published ref.

**Rationale:** Guarantees that the exact same commit SHA that was reviewed/drafted is the one being published.

**Status:** Accepted.

## ADR-004: Client-Side Encryption for Assets

**Context:** Git gateways or mirror repositories may be untrusted.

**Decision:** Encrypt all binary assets (images, PDFs) client-side using AES-256-GCM before uploading.

**Rationale:** defense-in-depth; the gateway only ever receives opaque encrypted blobs and an authenticated manifest.

**Status:** Accepted.

# 10 Quality Requirements

## 10.1 Quality Tree

The primary quality attributes for Git CMS are prioritized as follows:

1. **Security:** Cryptographic integrity and asset confidentiality.
2. **Simplicity:** Zero external database dependencies.
3. **Auditability:** Full provenance via Git's Merkle DAG.
4. **Performance:** Sub-second reads for standard blog workloads.

## 10.2 Quality Scenarios

### 10.2.1 QS-1: Tamper Detection

**Scenario:** An attacker modifies a published article directly on the Git gateway.

**Stimulus:** Malicious rewrite of Git history (`filter-branch`).

**Response:** The Merkle DAG checksum mismatch is immediately detected by any client pulling the update.

**Metric:** 100% detection of unauthorized history rewrites.

### 10.2.2 QS-2: Confidentiality

**Scenario:** A repository mirror is compromised.

**Stimulus:** Attacker attempts to view private image assets.

**Response:** Only AES-256-GCM ciphertext is visible; plaintext remains unrecoverable without the client-side key.

**Metric:** 0% leakage of plaintext assets.

## 11 Risks & Technical Debt

### 11.1 Risk 1: SHA-1 Collision

Git's reliance on SHA-1 is a known cryptographic risk. While the likelihood of a practical attack on a blog is low, the system should monitor Git's transition to SHA-256.

### 11.2 Risk 2: Repository Growth

Every draft save creates a permanent commit. Over years of active use, the object store could grow significantly. Regular `git gc -aggressive` and ref pruning strategies are needed.

### 11.3 Technical Debt Summary

Item	Priority	Impact
Automated ref pruning	High	Reduces repo bloat
Retry logic for CAS conflicts	Medium	Improves concurrent editing
External index for large ref counts	Medium	Improves <code>listArticles</code> performance

## 12 Glossary

**AES-256-GCM** Advanced Encryption Standard with 256-bit keys in Galois/Counter Mode.

**Bare Repository** A Git repository without a working directory, typically used on servers.

**CAS** Content-Addressable Store (or Compare-and-Swap, depending on context).

**Commit** A snapshot of the repository at a point in time.

**Empty Tree** The unique OID (`4b825dc...`) of a tree containing zero files.

**Merkle DAG** A directed acyclic graph where each node is identified by the hash of its content.

**Ref** A pointer to a Git object (e.g., branch, tag, or article slug).

**Trailer** RFC 822 metadata at the end of a commit message.

## A Appendix A: Example Commands

### A.1 Draft an Article

```
1 echo "# My First Post" | git cms draft hello-world "My First Post"
```

## A.2 Publish an Article

```
1 git cms publish hello-world
```

## A.3 List All Drafts

```
1 git cms list
```

## A.4 Upload Asset

```
1 git cms upload hello-world image.png
```

## B Appendix B: Directory Structure

```
1 git-cms/
2   +- bin/
3   |   +- git-cms.js          # CLI entry point
4   +- src/
5   |   +- lib/
6   |   |   +- CmsService.js    # Core orchestrator
7   |   +- server/
8   |   |   +- index.js        # HTTP API server
9   +- test/
10  |   +- git.test.js         # Integration tests
11  |   +- e2e/                # Playwright tests
12  +- public/                # Static admin UI
13  +- docs/                  # Documentation
```

## C Appendix C: Related Projects

- **git-stargate:** Git gateway for enforcingFF-only and signing.
- **git-stunts:** Lego blocks for Git plumbing.

## D Appendix D: References

1. Git Internals (Pro Git Book)
2. RFC 822 (Internet Message Format)
3. AES-GCM (NIST SP 800-38D)

## E Appendix D: References

1. **Git Internals (Pro Git Book):**  
<https://git-scm.com/book/en/v2/Git-Internals-Plumbing-and-Porcelain>
2. **RFC 822 (Internet Message Format):**  
<https://tools.ietf.org/html/rfc822>
3. **AES-GCM (NIST SP 800-38D):**  
<https://csrc.nist.gov/publications/detail/sp/800-38d/final>